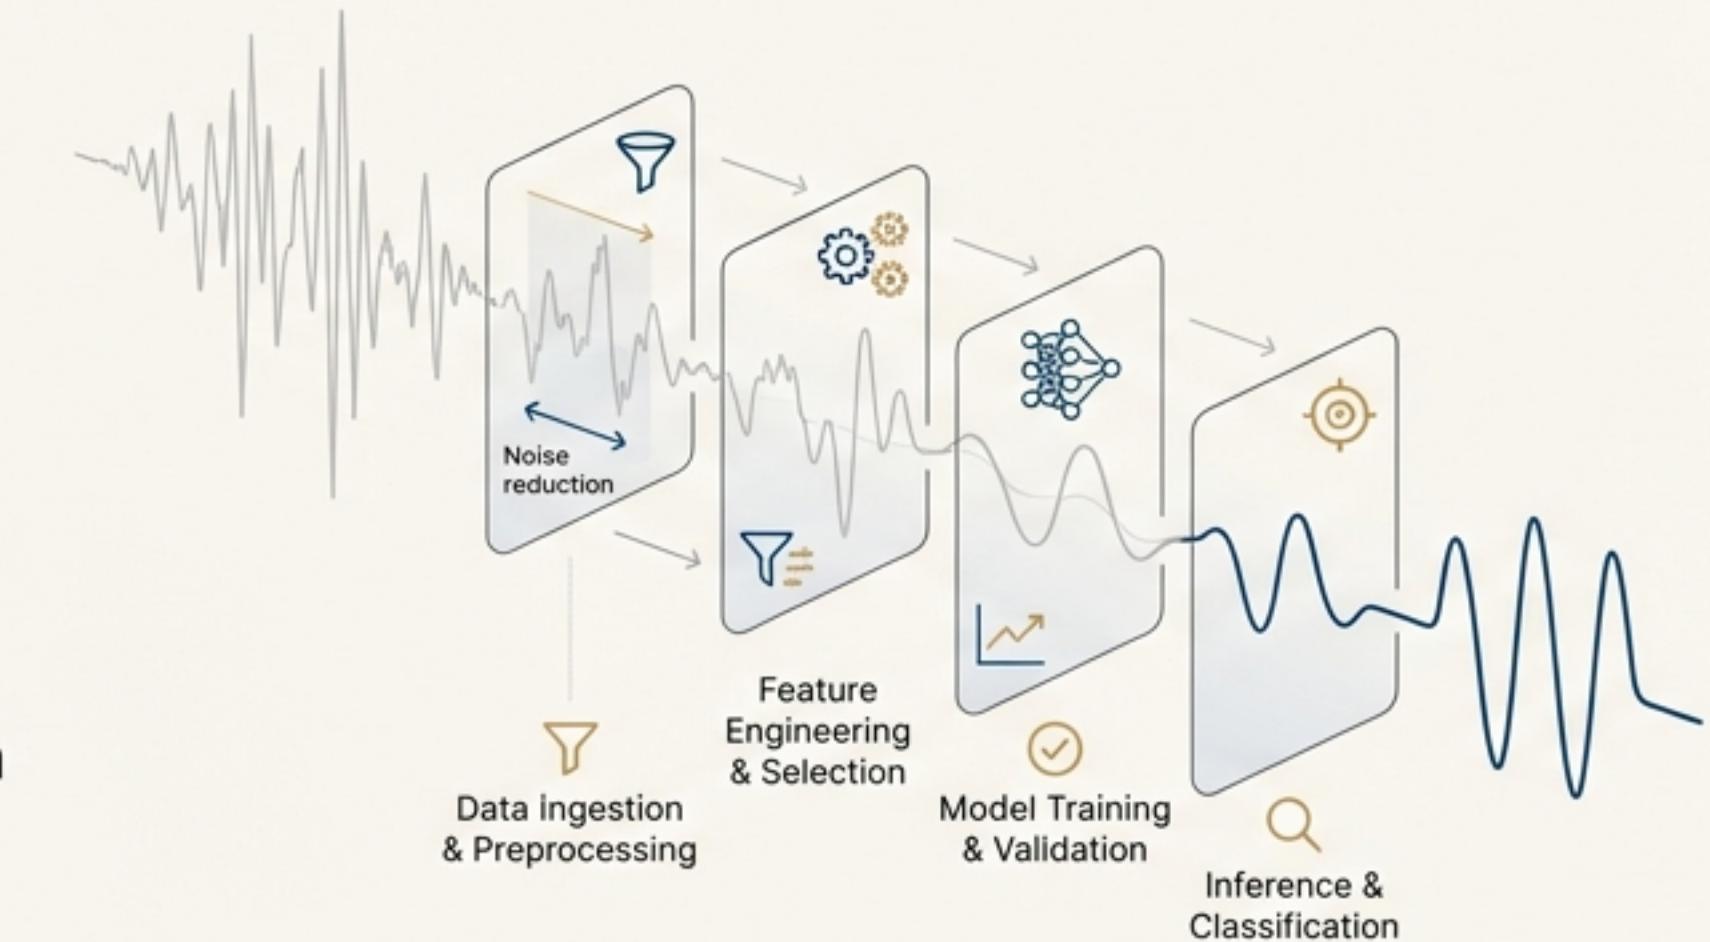
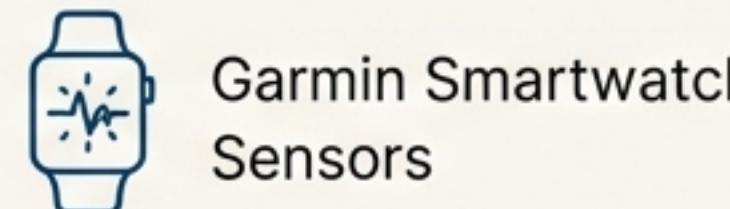


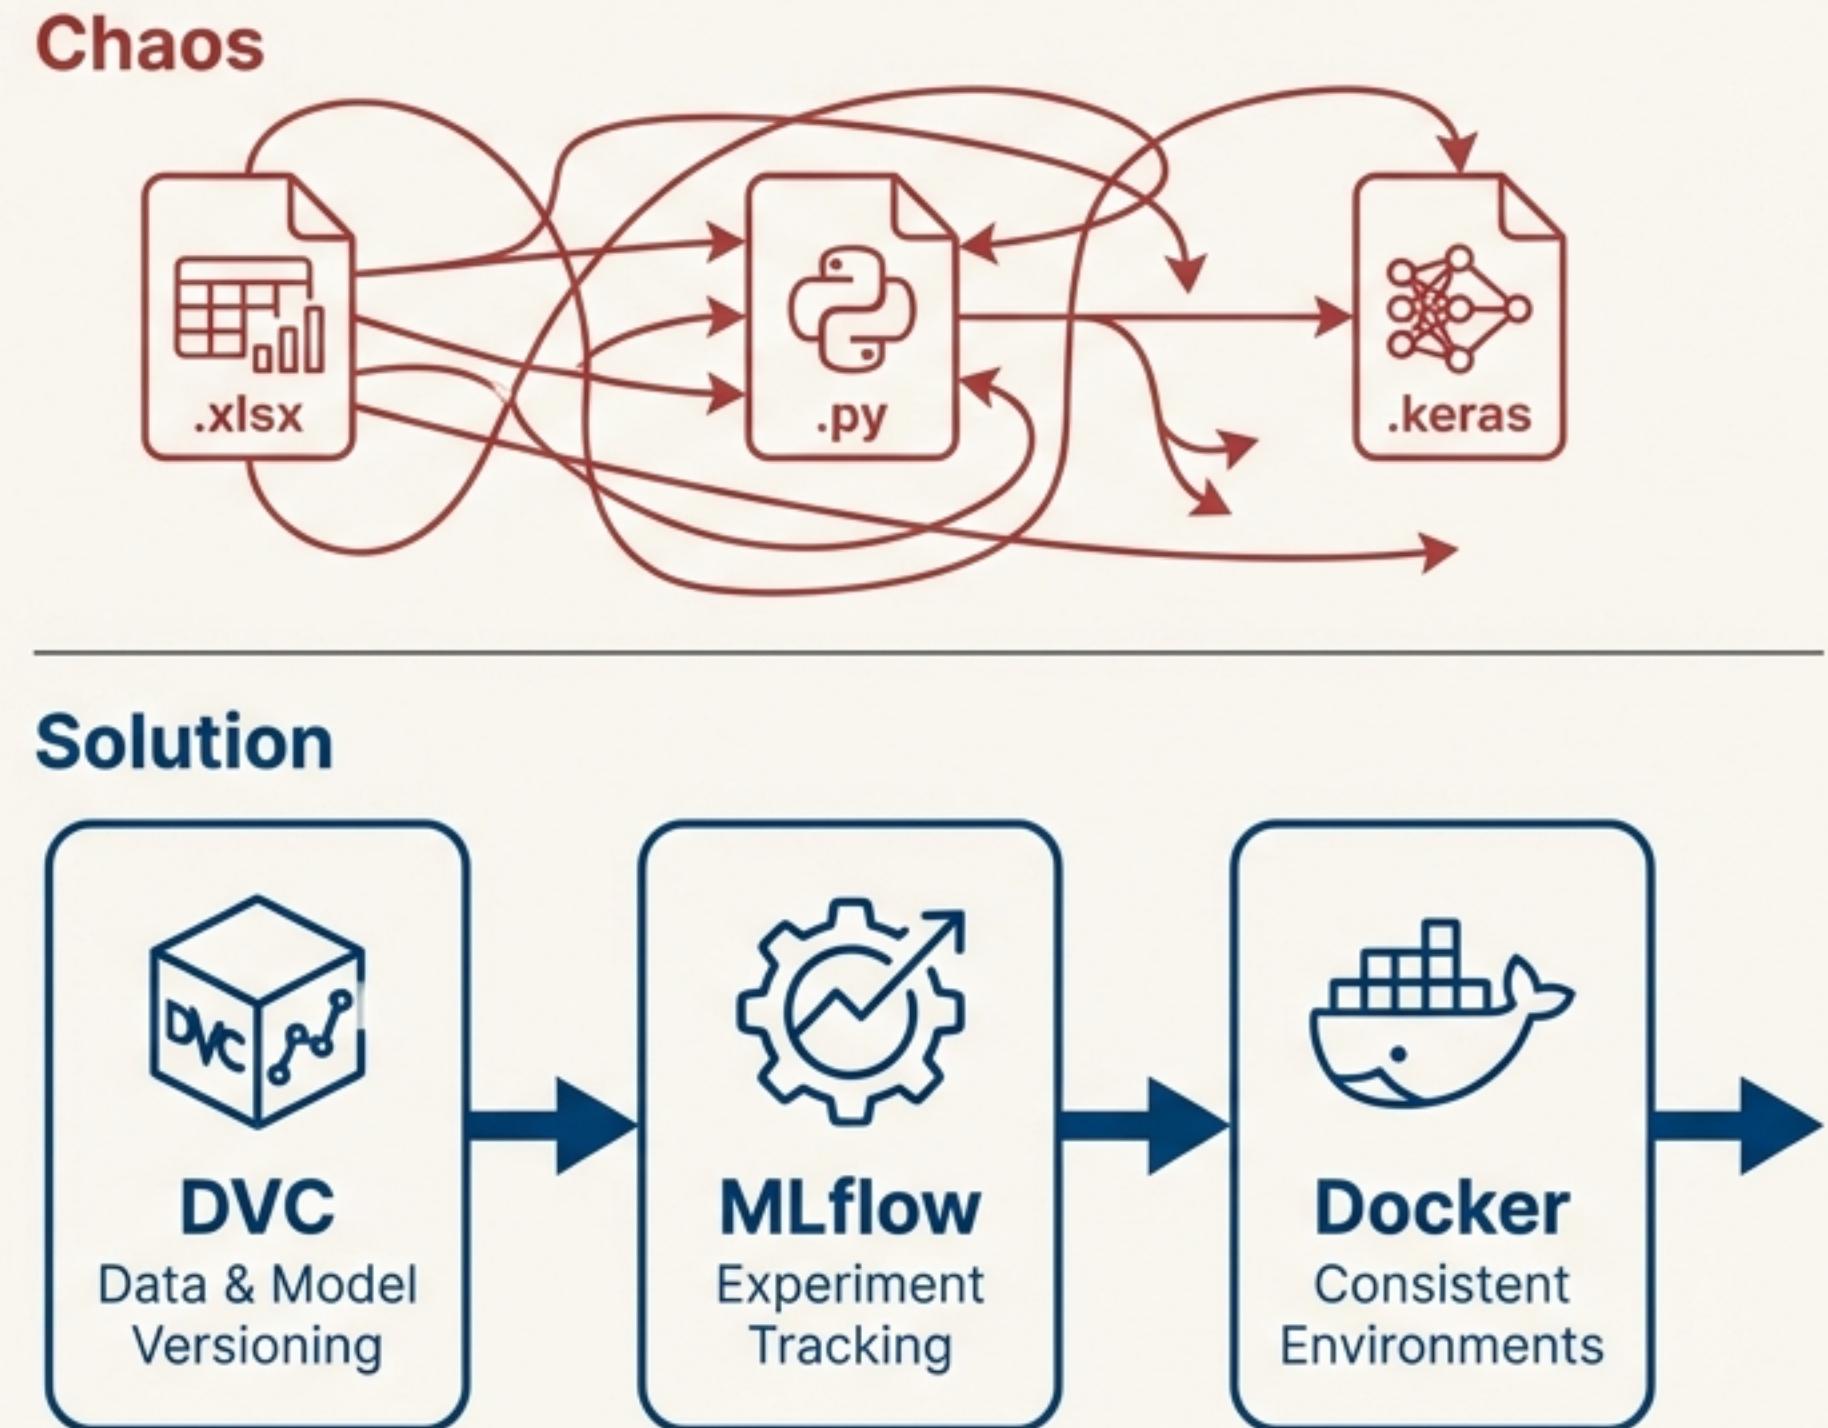
Deconstructing the MLOps Pipeline for Human Activity Recognition

A technical walkthrough of the end-to-end system for recognizing anxiety-related activities from wearable sensor data.

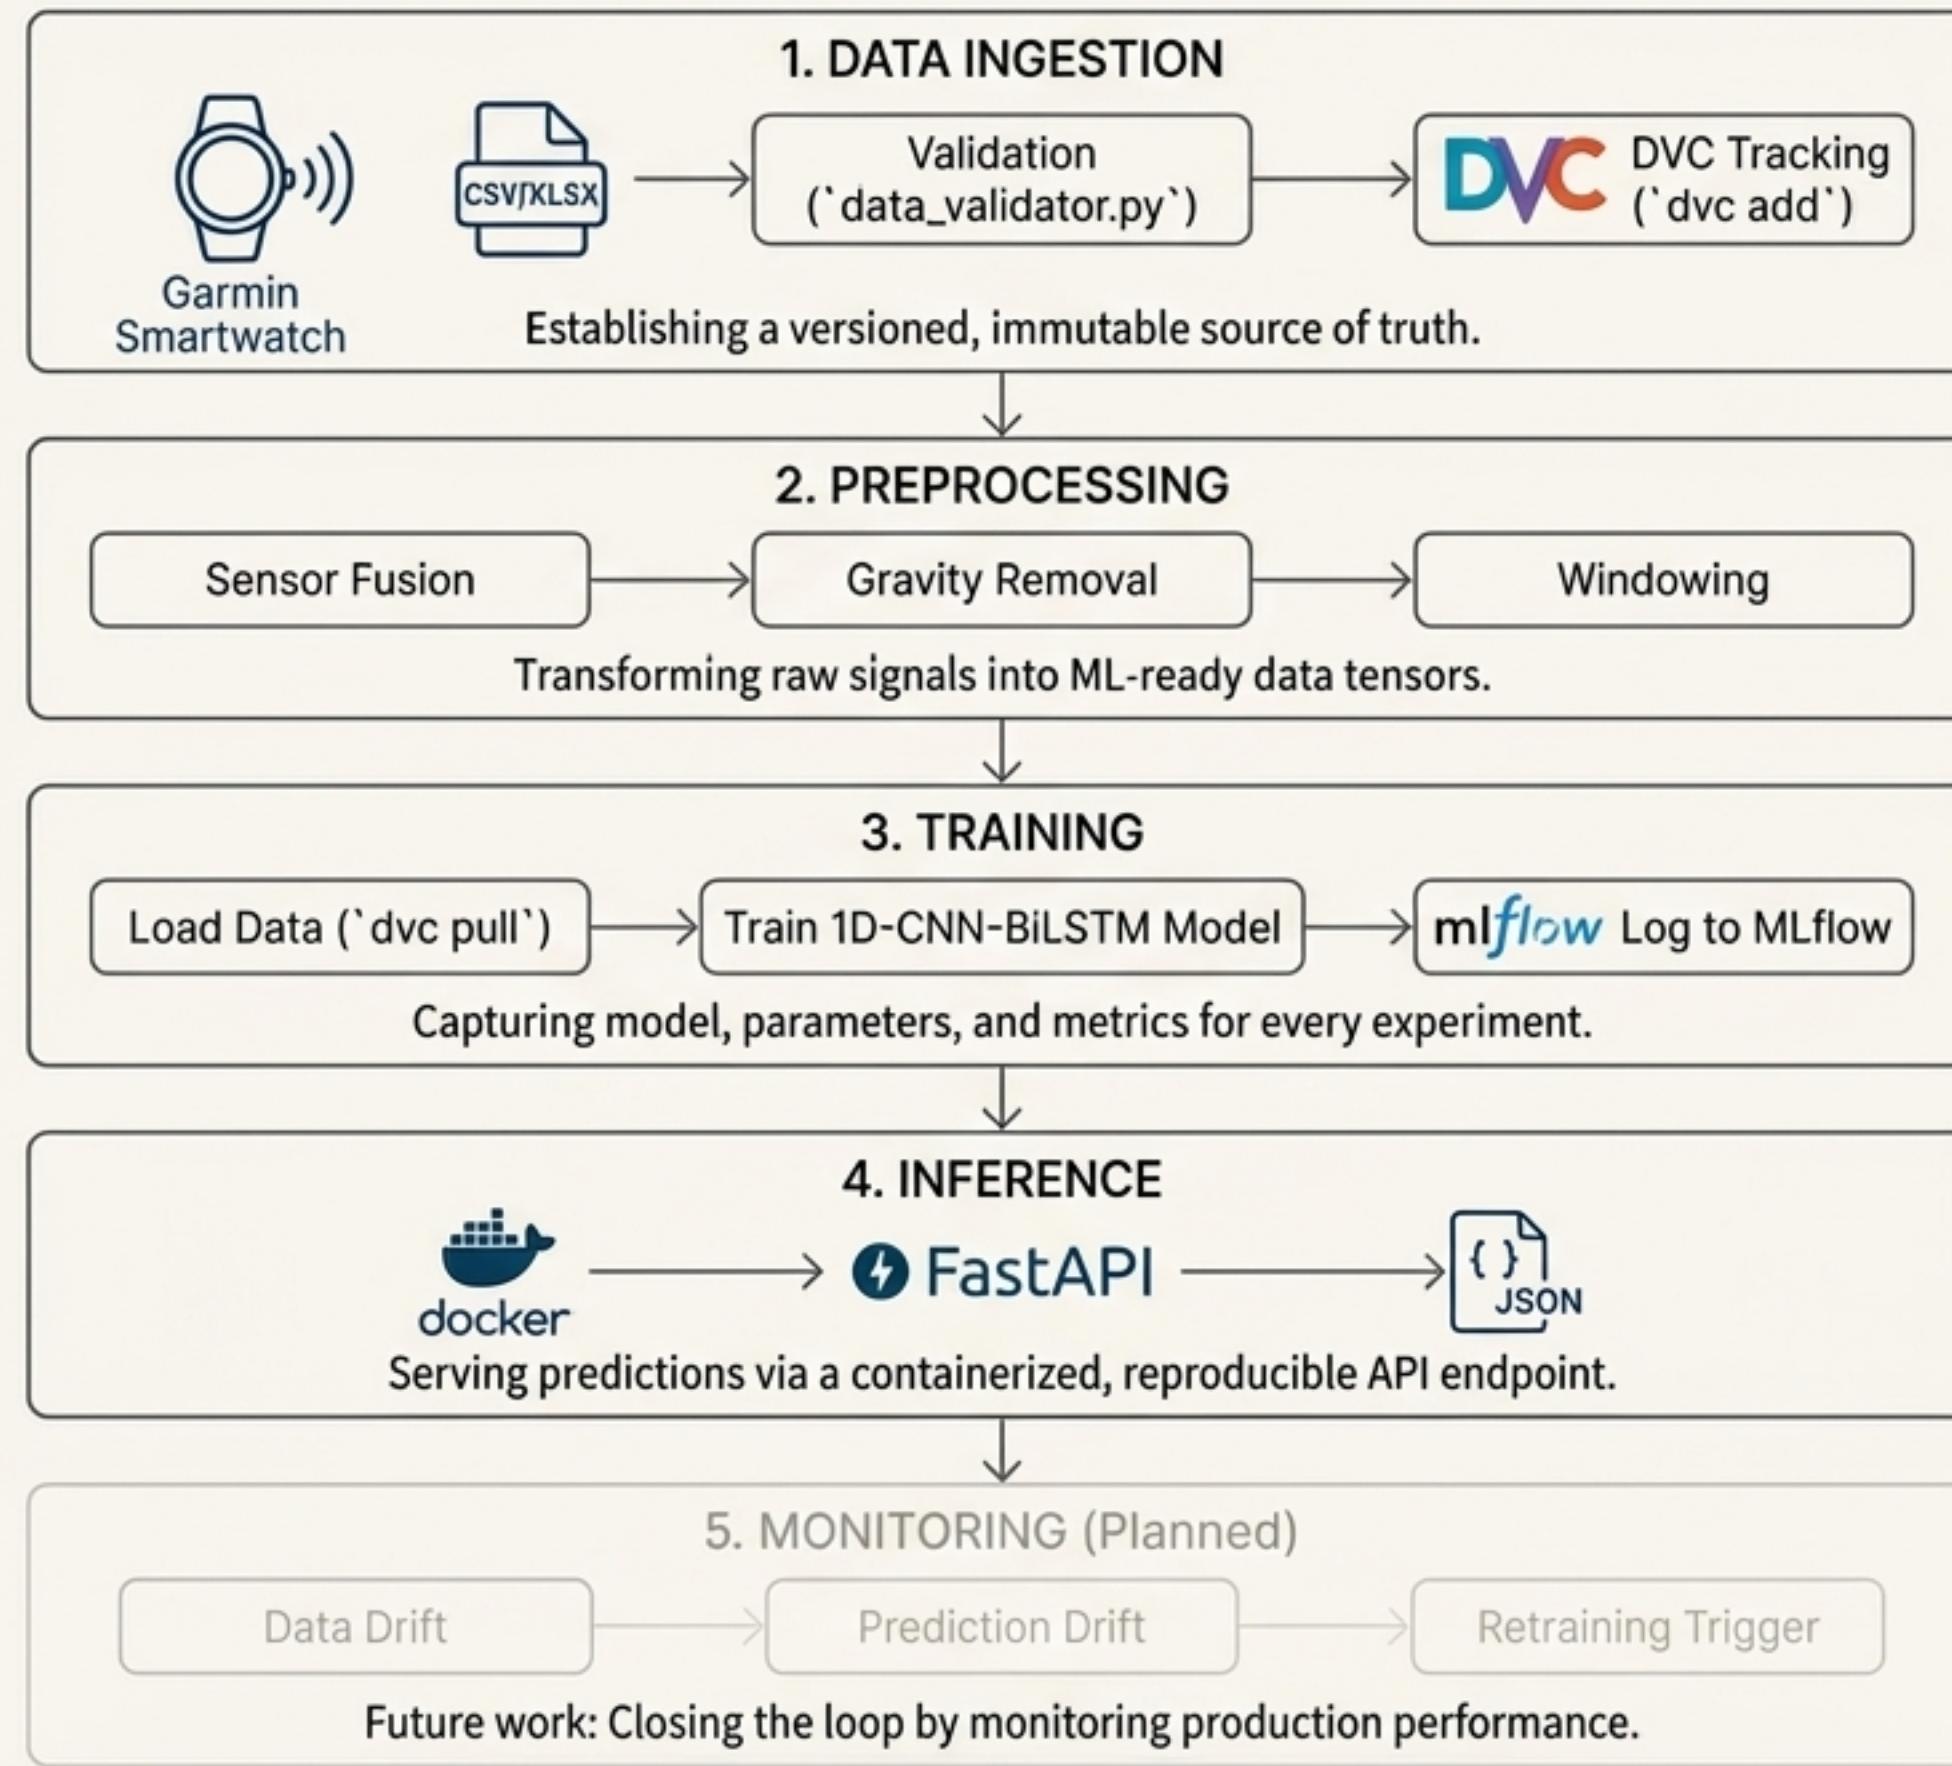


The MLOps Imperative: From Data Chaos to Reproducible Insight

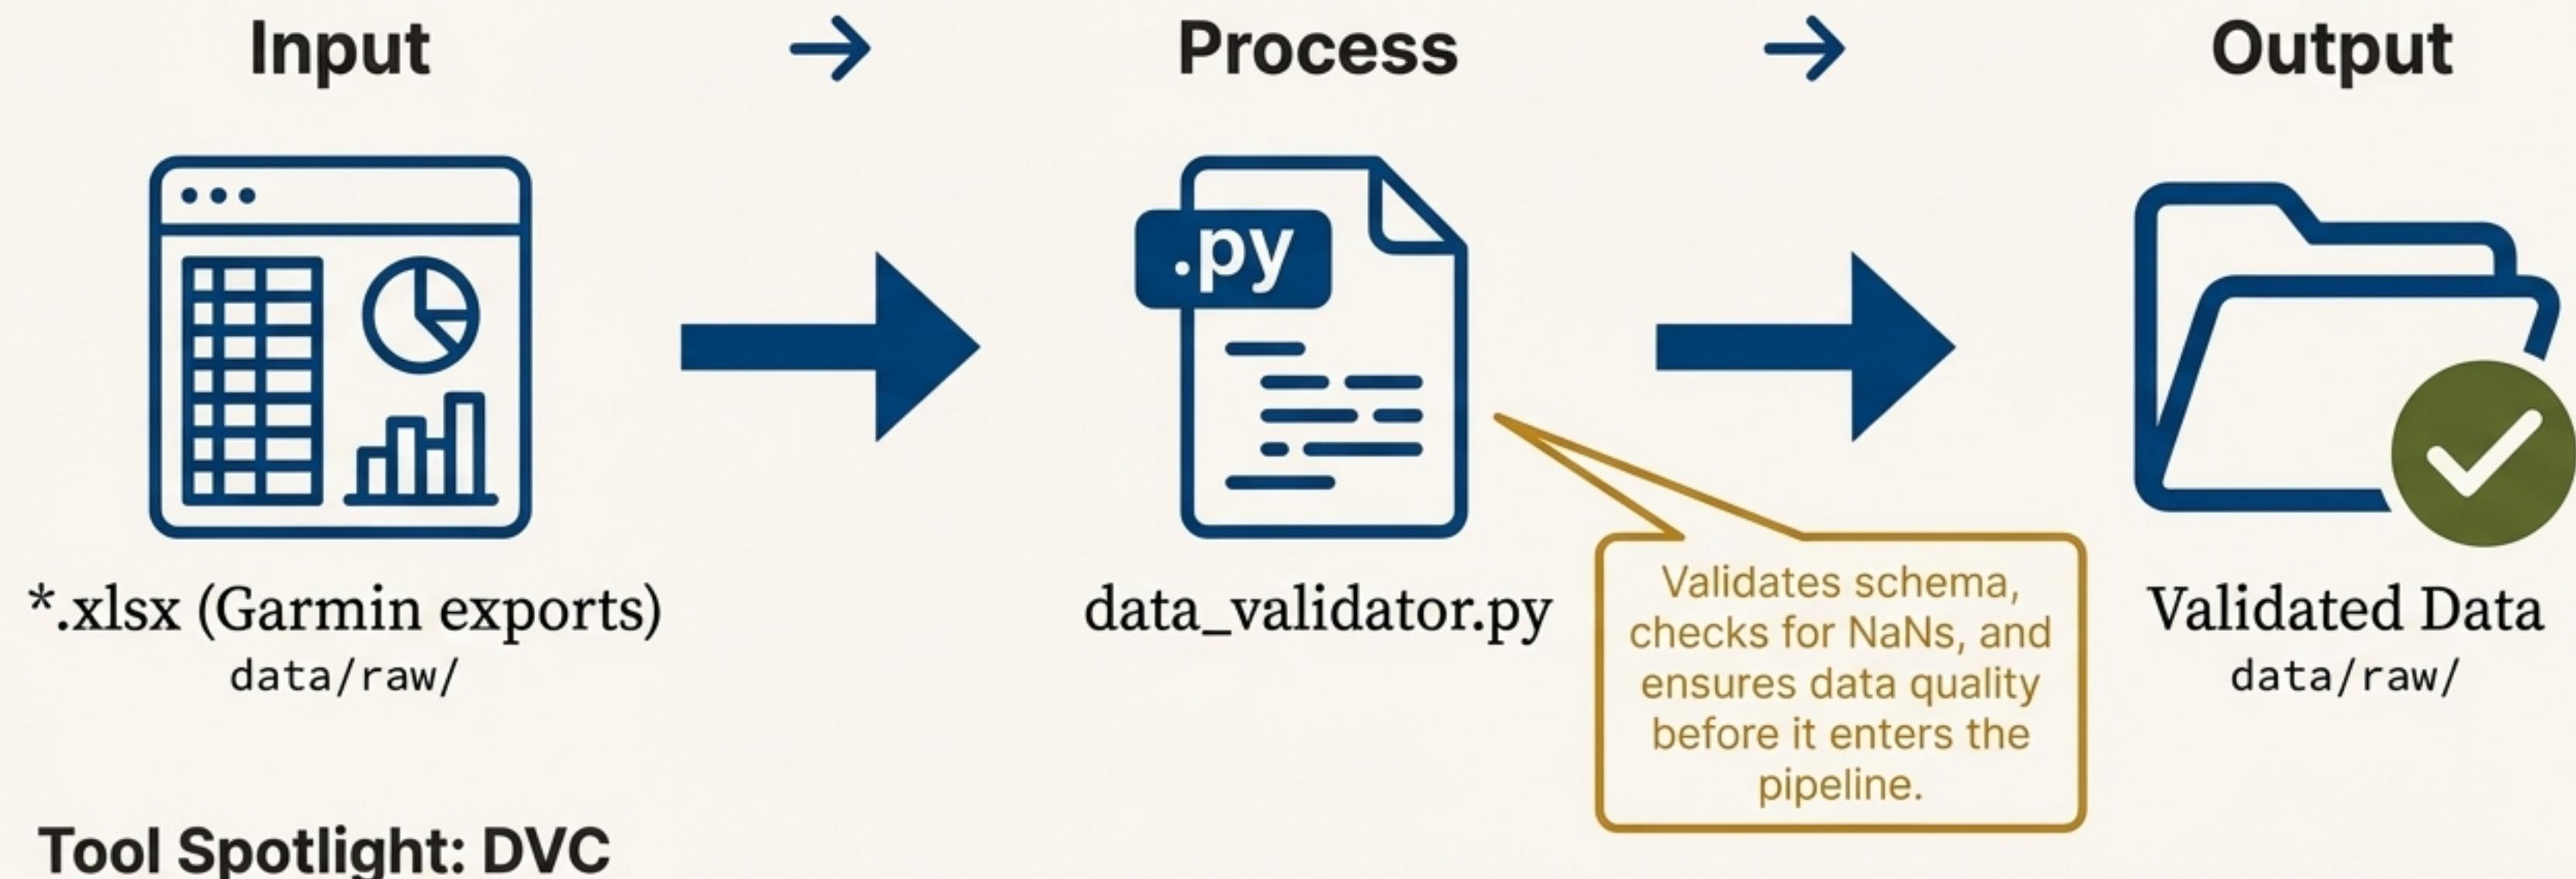
- **Data & Model Versioning:** How do we track massive datasets and the models trained on them without cluttering our Git history?
- **Experiment Traceability:** How can we reliably reproduce results and compare which training parameters led to the best model?
- **Deployment Consistency:** How do we eliminate the “it works on my machine” problem and ensure our model performs identically in production?



The End-to-End Blueprint



Stage 1: Ingestion & Versioning the Source of Truth



Tool Spotlight: DVC



```
dvc add data/raw
```

DVC takes over tracking for the large XLSX files. Git now only stores a tiny .dvc pointer file, keeping the repository lightweight.

Stage 2: Transforming Raw Signals into ML-Ready Tensors



Part 1: Sensor Fusion & Cleaning

`sensor_data_pipeline.py`

- Fuses accelerometer & gyro
- Resamples to 50Hz
- Removes gravity component (0.3Hz HPF)

data/processed/sensor_fused_50Hz.csv

Part 2: Data Preparation

`preprocess_data.py`

- Creates sliding windows (200 timesteps)
- Applies StandardScaler normalization

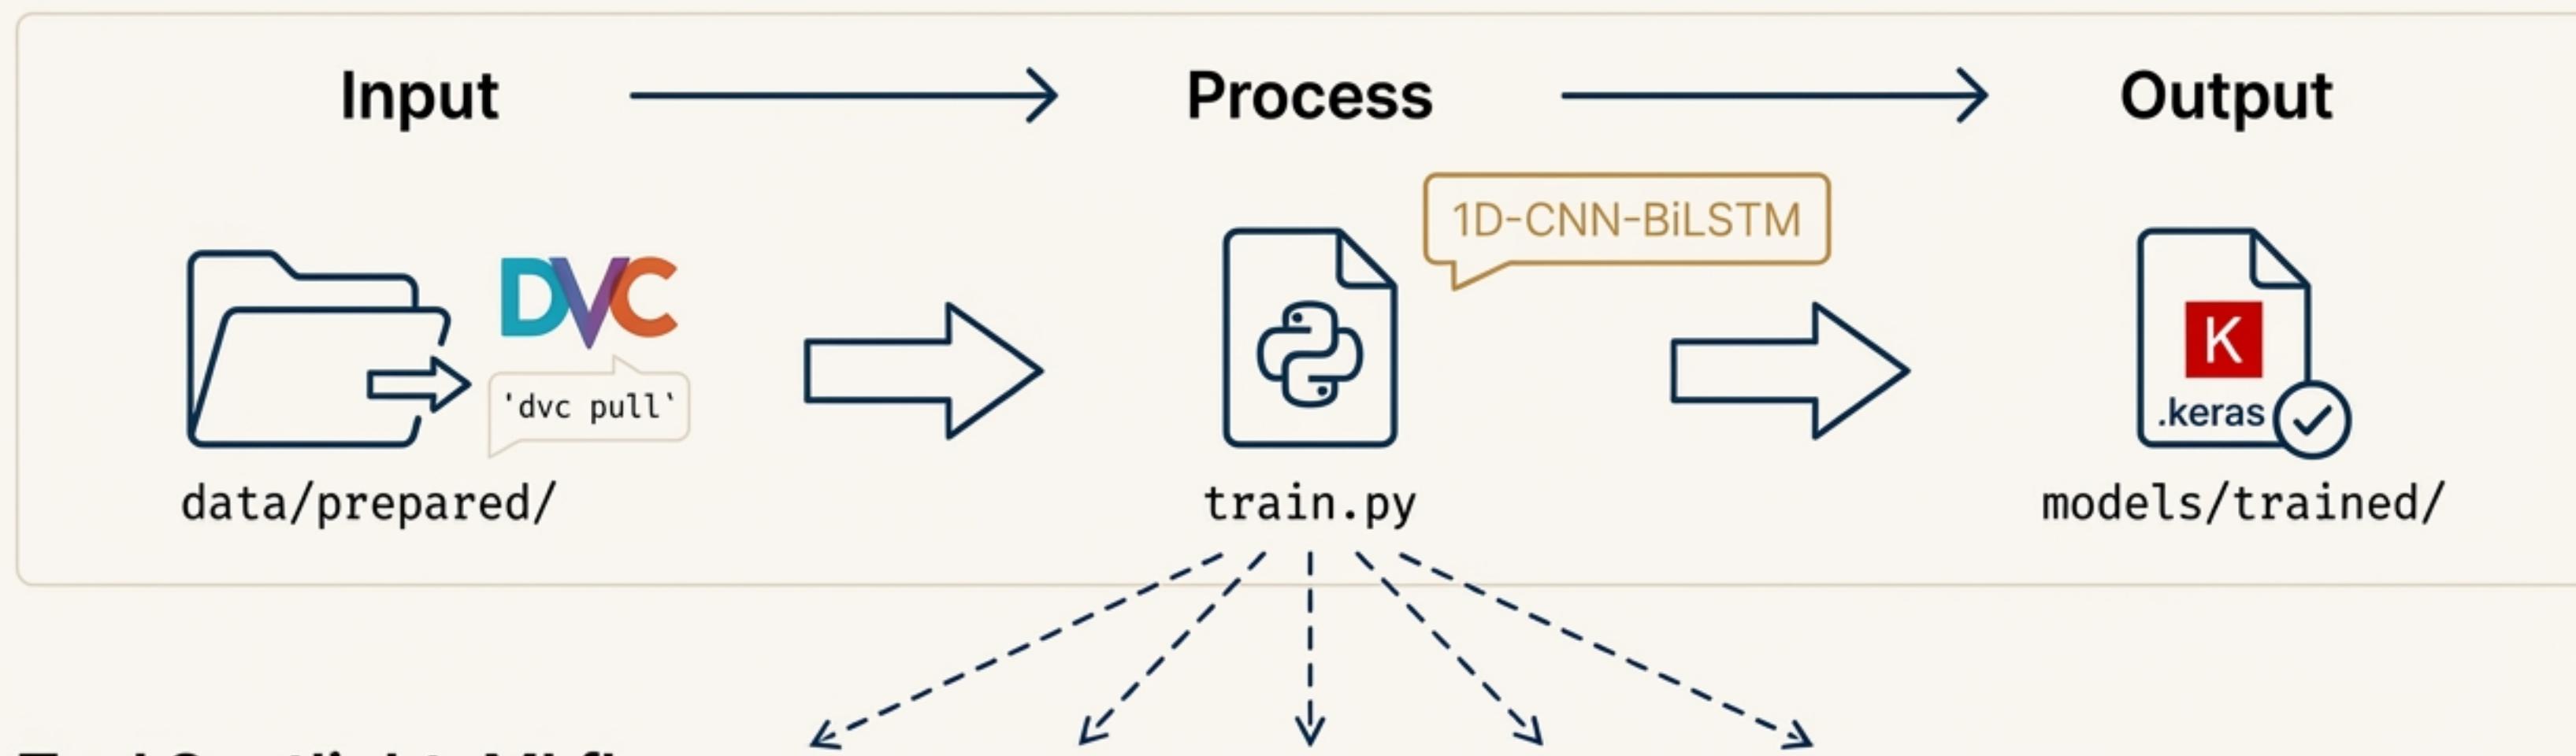


data/prepared/train_X.npy



data/prepared/train_y.npy

Stage 3: Training the Model with Full Experimental Visibility



Tool Spotlight: MLflow

mlflow

The Experiment Hub: Comparing Runs and Promoting Models

Run Comparison

Run Name	learning_rate	window_size	accuracy	f1_macro
training_v1	0.001	200	0.95	0.93
training_v2_lr_halved	0.0005	200	0.92	0.90

Model Registry

har-1dcnn-bilstm v1

Register Model

The best performing model is promoted to the Model Registry for staging and production.

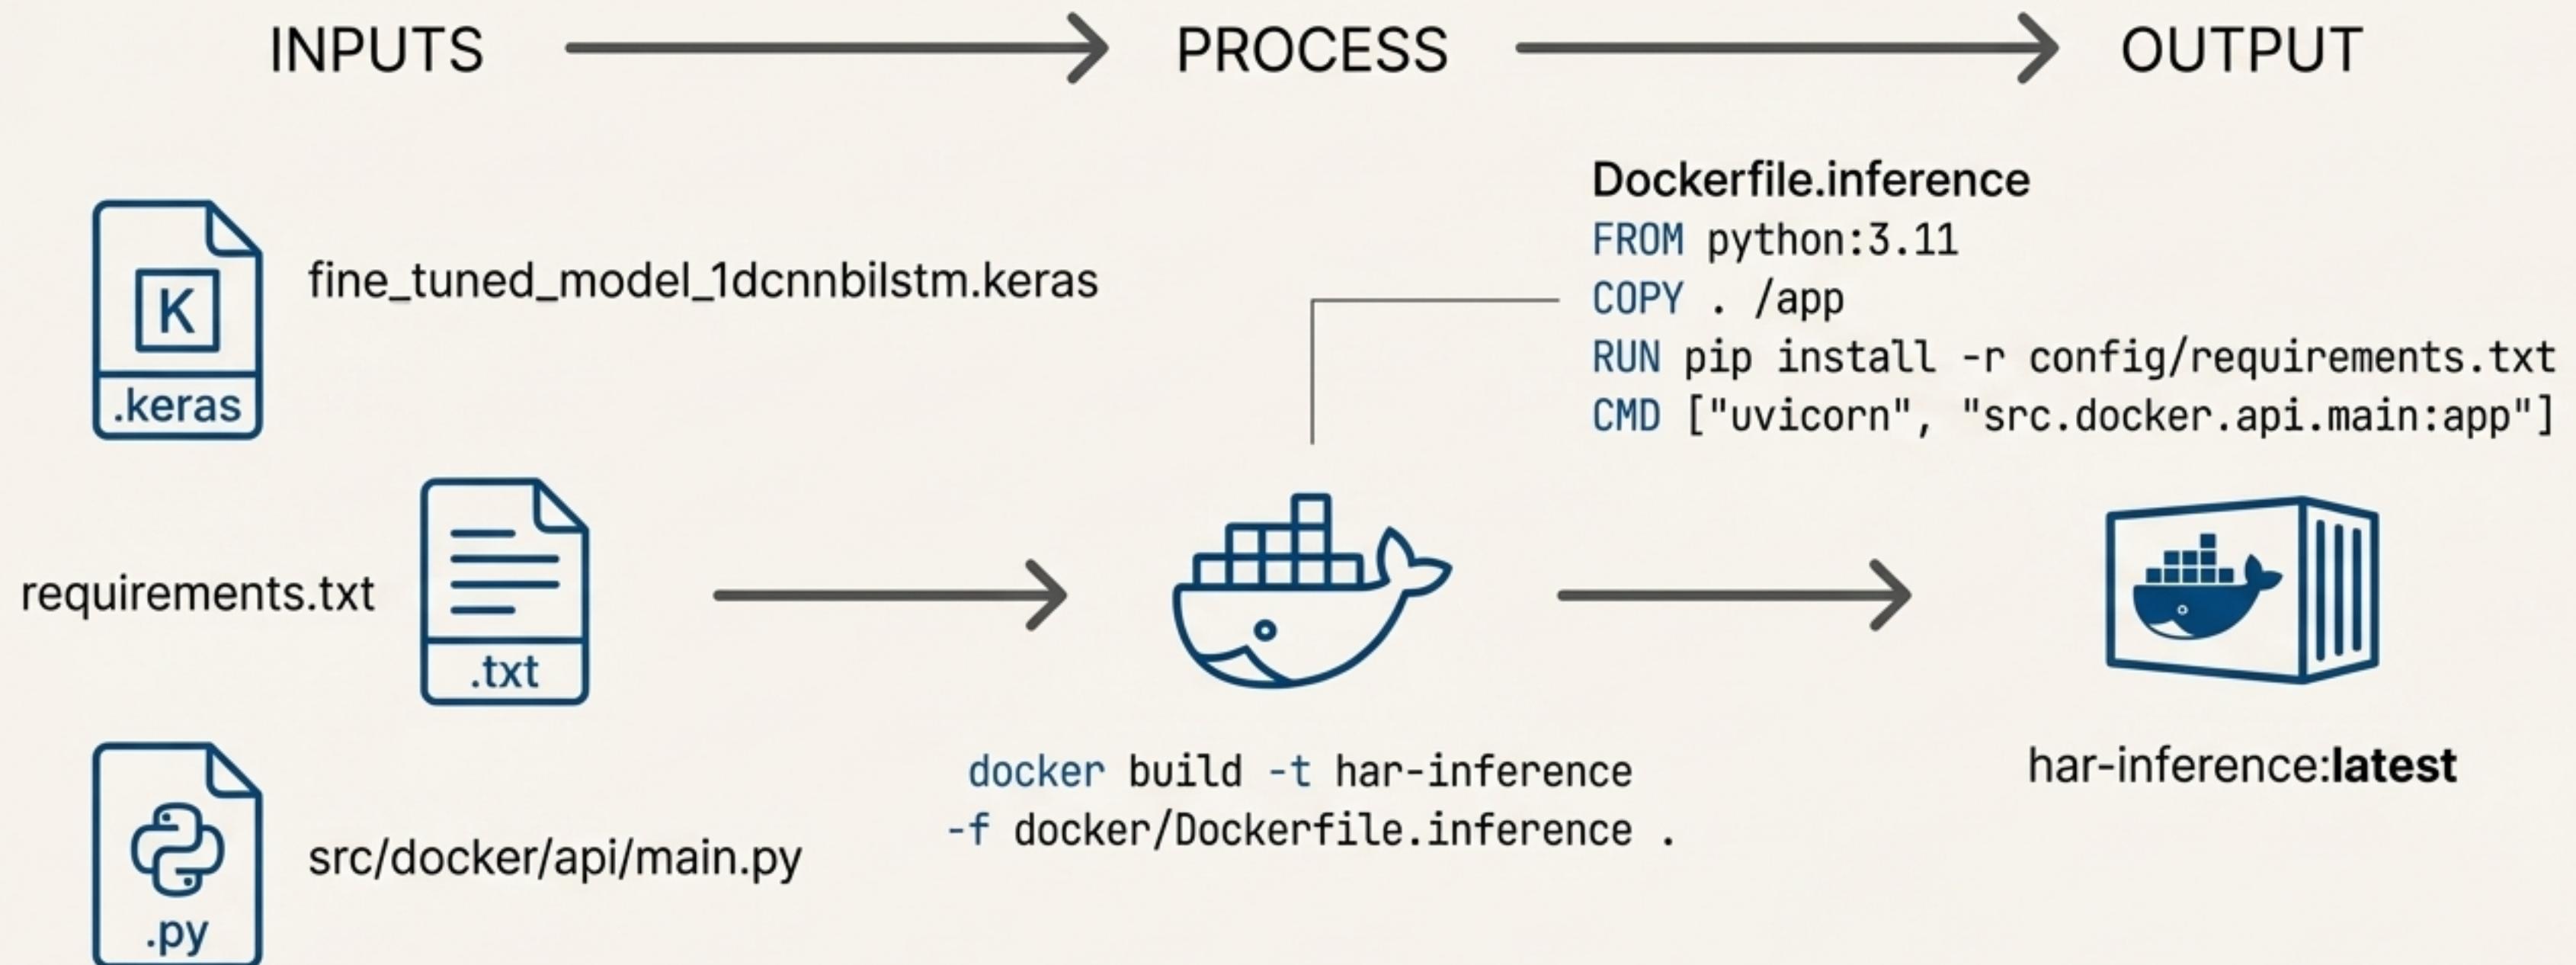
Artifact Visualization

		Walking	Sitting	Standing	Laying
Actual Class	Predicted Class	Walking	Sitting	Standing	Laying
		Walking	True Positives	Misclassifications	Misclassifications
Sitting	Misclassifications	True Positives	Misclassifications	Misclassifications	
Standing	Misclassifications	Misclassifications	True Positives	Misclassifications	
Laying	Misclassifications	Misclassifications	Misclassifications	True Positives	

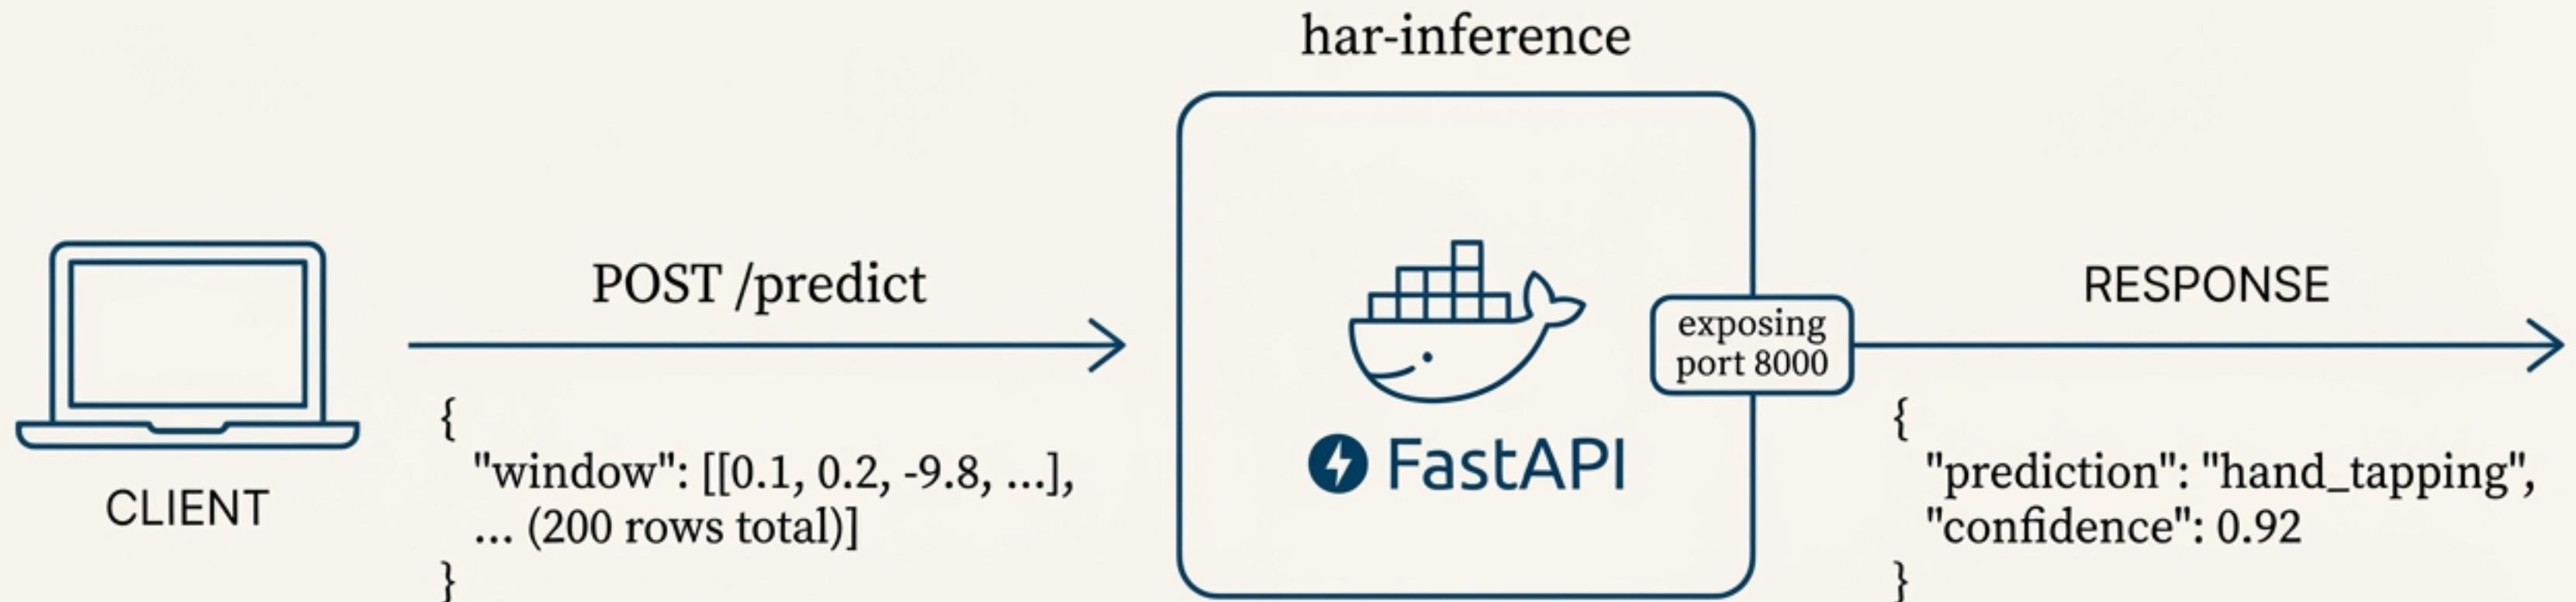
Legend: █ True Positives █ Misclassifications

> Accessible locally at <http://localhost:5000> via Docker Compose.

Stage 4, Part 1: Packaging the Predictor for Production



Stage 4, Part 2: Serving Live Predictions on Demand



Tool Spotlight: Docker Compose



`docker-compose.yml`

The `'docker-compose.yml'` file orchestrates the entire service stack, including the `'inference'` API and the `'mlflow'` server.

A Key Finding: Solving the Domain Shift Catastrophe

Initial deployment resulted in the model predicting "hand_tapping" for 95% of all production data, despite high validation accuracy.

Root Cause Analysis

Training Data

Gravity component was **removed**. Mean Az acceleration: $\sim 0 \text{ m/s}^2$.

Production Data

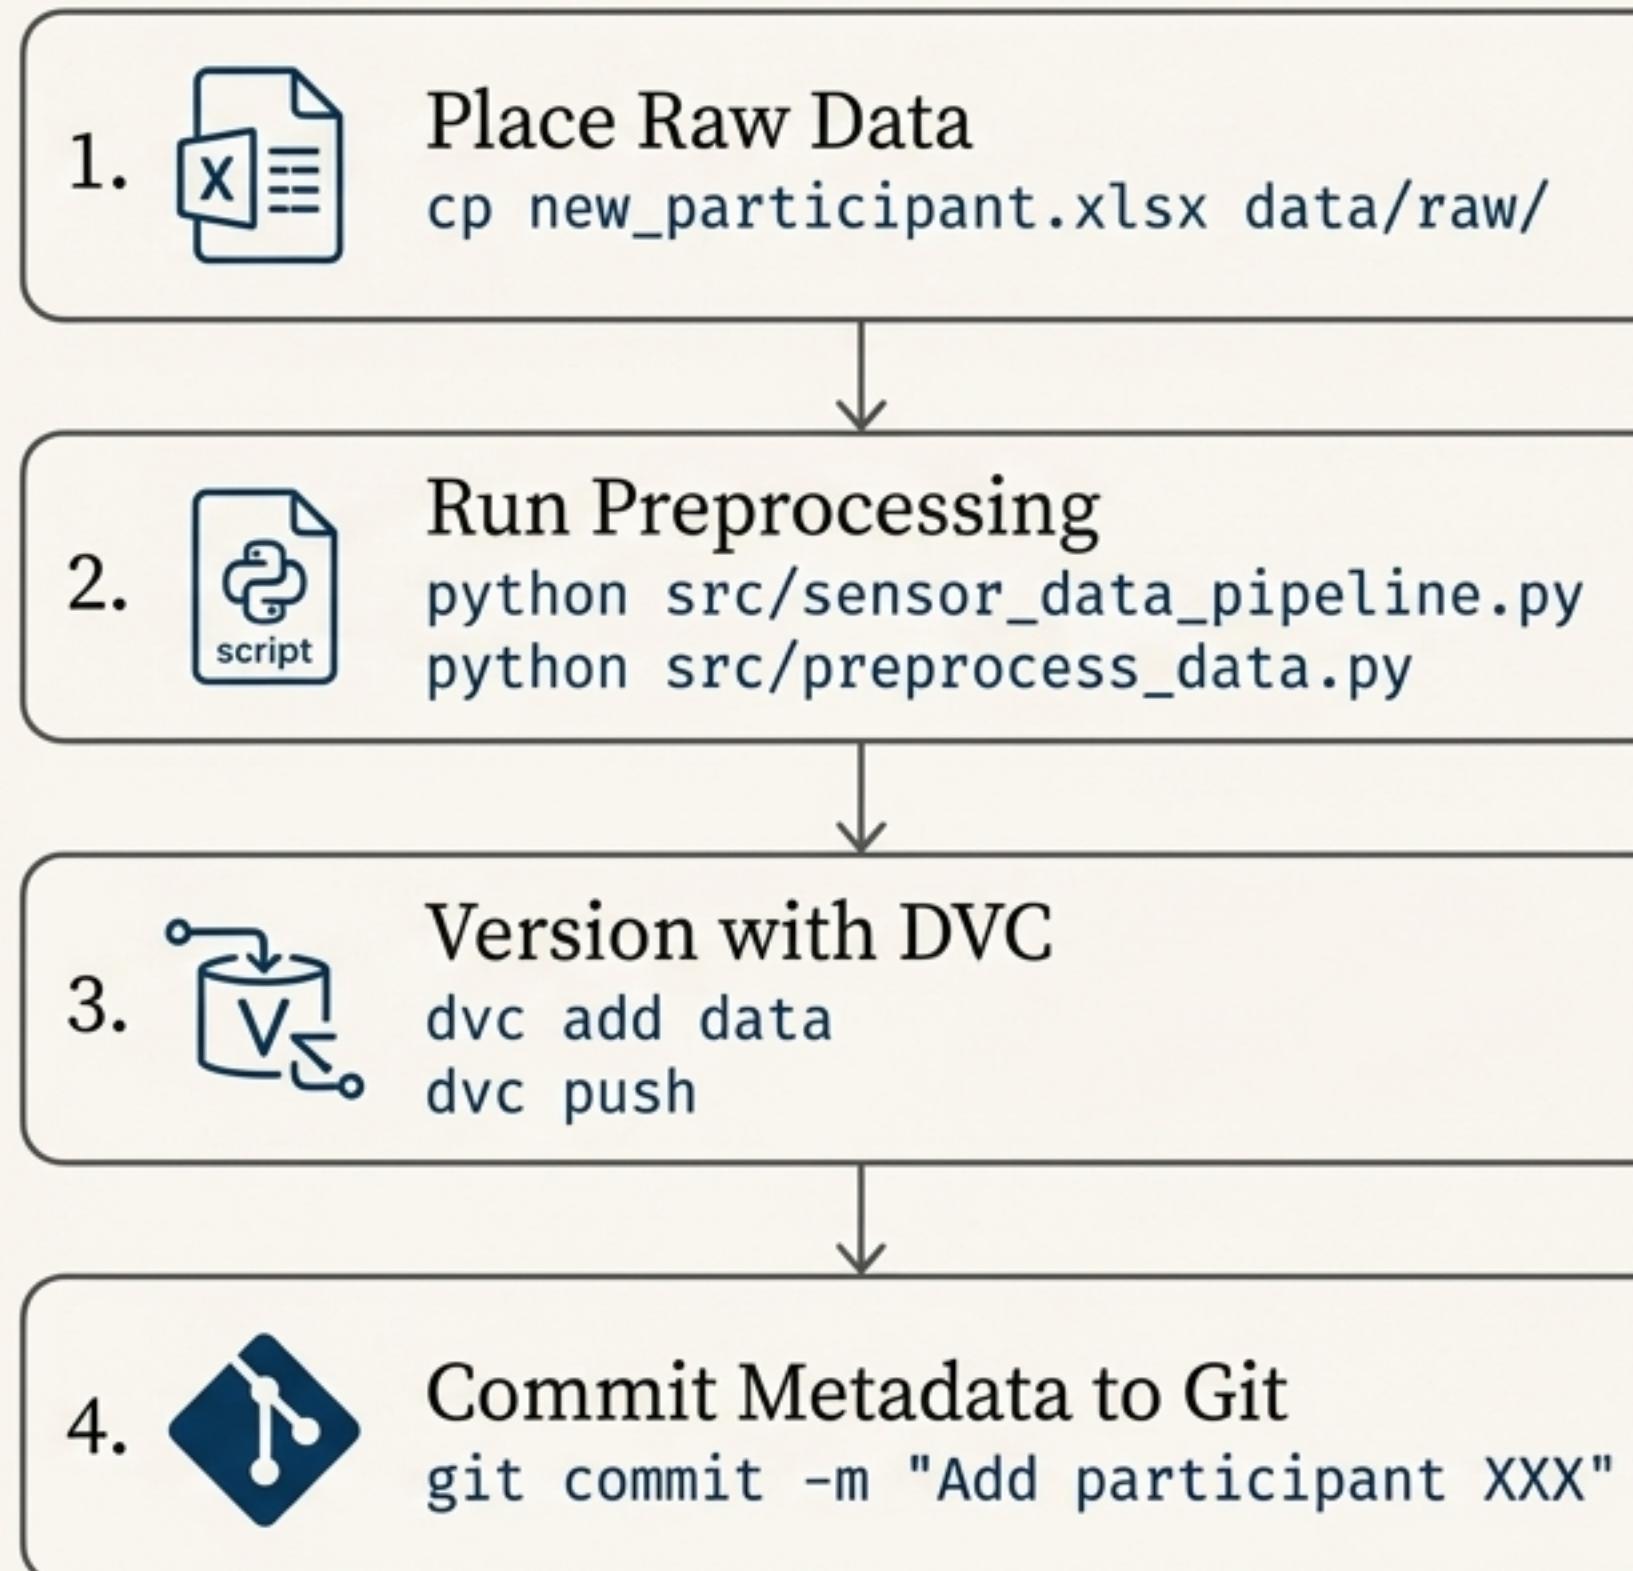
Gravity component was **present**. Mean Az acceleration: -9.83 m/s^2 .

Before vs. After

Metric	Before Fix	After Fix
hand_tapping %	95.4%	4.2%
Unique Classes Predicted	4 / 11	7 / 11

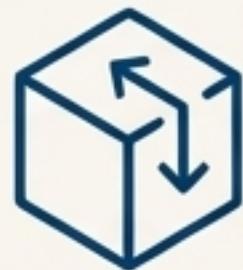
Implemented a configurable Butterworth high-pass filter (0.3 Hz) in the preprocessing pipeline (`sensor_data_pipeline.py`) to remove the gravity vector from all incoming data, controlled via `config/pipeline_config.yaml`.

The Workflow in Action: Onboarding New Participant Data



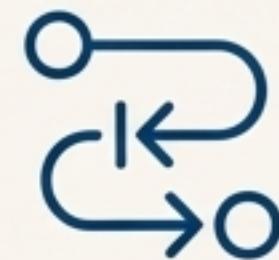
This standardized workflow ensures all new data is processed, validated, and versioned consistently before it is used for inference or retraining.

A Blueprint for Reproducible Machine Learning



Data & Model Versioning

Guarantees data integrity and full pipeline reproducibility. Every dataset and artifact is versioned and auditable.



Experiment Tracking & Governance

Provides complete experimental clarity. Every training run is logged, making results easy to compare, reproduce, and promote.



Environment Containerization & Deployment

Delivers production consistency. Every deployment is identical, eliminating environmental drift and the “it works on my machine” problem.

Explore the Project



GitHub Repository: The complete source code, documentation, and setup instructions are available on GitHub.

``ShalinVachheta017/MasterArbeit_MLops-``



API Documentation: Once running, interactive API documentation is available via the Swagger UI at ``http://localhost:8000/docs``.

