

THEESIS_PROGRESS_ANALYSIS

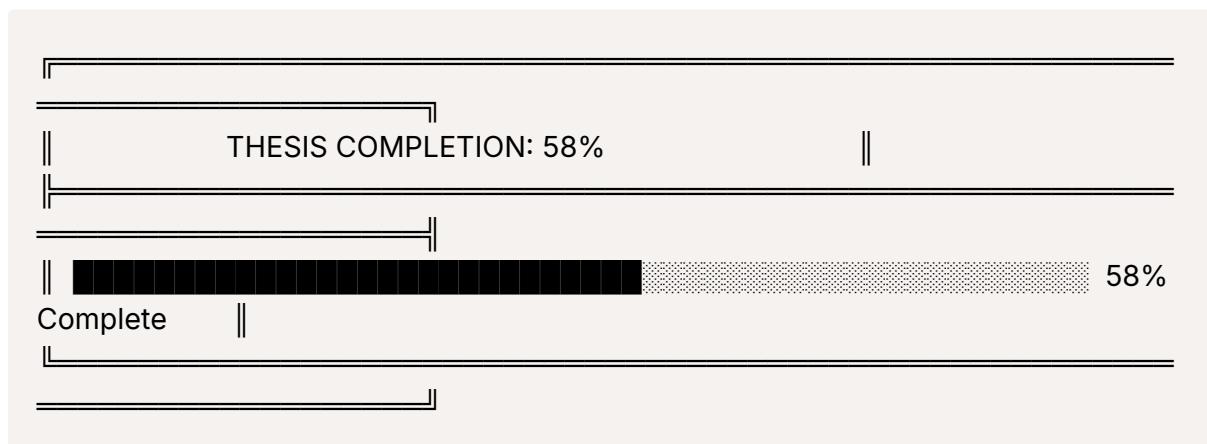
Thesis Progress Analysis & Improvement Roadmap

Thesis: Developing a MLOps Pipeline for Continuous Mental Health Monitoring using Wearable Sensor Data

Analysis Date: December 13, 2025

Timeline: October 2025 - April 2026 (6 months)

Overall Completion Status



Phase	Status	Completion
Month 1: Data Ingestion	✓ Complete	100%
Month 2: Model Training & Versioning	✓ Complete	95%
Month 3: CI/CD & Deployment	🟡 Partial	60%
Month 4: Integration & Monitoring	🟡 Partial	40%
Month 5: Refinement & Retraining	🔴 Not Started	10%
Month 6: Thesis Writing	🔴 Not Started	5%

Progress Visualization (Mermaid)

```
gantt
    title Thesis Progress - MLOps Pipeline for Mental Health Monitoring
    dateFormat YYYY-MM-DD
```

```
section Month 1: Data
```

Literature Review :done, m1w1, 2025-10-01, 14d
Data Ingestion Pipeline :done, m1w3, 2025-10-15, 14d

section Month 2: Training

Automated Training Loop :done, m2w1, 2025-11-01, 14d
MLflow Experiment Tracking :done, m2w3, 2025-11-15, 14d
Docker Containerization :done, m2w4, 2025-11-22, 7d

section Month 3: CI/CD

Docker Compose Setup :done, m3w1, 2025-12-01, 7d
FastAPI Inference API :done, m3w2, 2025-12-08, 7d
GitHub Actions CI/CD :active, m3w3, 2025-12-15, 14d

section Month 4: Monitoring

Data Drift Detection :active, m4w1, 2025-12-29, 14d
Prediction Drift :m4w2, 2026-01-12, 14d
Performance Monitoring :m4w3, 2026-01-26, 7d

section Month 5: Refinement

Pipeline Robustness :m5w1, 2026-02-01, 14d
Retraining Strategy :m5w2, 2026-02-15, 14d

section Month 6: Writing

Thesis Document :m6w1, 2026-03-01, 21d
Final Presentation :m6w2, 2026-03-22, 7d



Detailed Task Completion Analysis

✓ Month 1: Research, Planning & Data Ingestion (100% Complete)

Task	Status	Evidence
Literature Review	✓ Done	research_papers/COMPREHENSIVE_RESEARCH_PAPERS_SUMMARY.md - 77+ papers analyzed
Understand 1D-CNN-BiLSTM	✓ Done	Model analyzed: 499,131 params, 11 classes
Define scope & metrics	✓ Done	Thesis_Plan.md , project documentation
Raw data ingestion	✓ Done	data/raw/*.xlsx (Garmin exports)
Preprocessing pipeline	✓ Done	src/sensor_data_pipeline.py , src/preprocess_data.py

Task	Status	Evidence
Automated & reproducible	✓ Done	Python scripts with config files

Deliverables:

- ✓ Sensor fusion (accelerometer + gyroscope)
- ✓ 50Hz resampling
- ✓ Domain calibration (-6.295 m/s² Az offset)
- ✓ 200-sample windows with 50% overlap

✓ Month 2: Model Training & Versioning (95% Complete)

Task	Status	Evidence
Automated training script	✓ Done	src/run_inference.py (896 lines)
MLflow logging	✓ Done	src/mlflow_tracking.py (654 lines)
Hyperparameter tracking	✓ Done	MLflow params logging
Model versioning	✓ Done	DVC + MLflow Model Registry
Experiment tracking	✓ Done	mlruns/ directory
Docker containerization	✓ Done	docker/Dockerfile.training , Dockerfile.inference

Deliverables:

- ✓ MLflow experiment tracking with metrics
- ✓ DVC for data/model versioning
- ✓ Docker containers for training and inference
- ⚡ Hyperparameter optimization (basic, not automated)

🟡 Month 3: CI/CD & Deployment (60% Complete)

Task	Status	Evidence
Docker Compose	✓ Done	docker-compose.yml (4 services)
FastAPI inference API	✓ Done	docker/api/main.py
Dockerfile.inference	✓ Done	docker/Dockerfile.inference
GitHub Actions CI/CD	✗ Missing	No .github/workflows/ directory
Automated testing	✗ Missing	tests/ folder empty
Docker registry push	✗ Missing	Not implemented

Deliverables:

- ✓ docker-compose.yml with MLflow + Inference services
- ✓ FastAPI /predict endpoint
- ✓ Containerized inference
- ✗ CI/CD pipeline (GitHub Actions)
- ✗ Automated unit tests

🟡 Month 4: Integration & Monitoring (40% Complete)

Task	Status	Evidence
Activity → Prognosis flow	🔴 Missing	Prognosis model not integrated
Data drift detection	🟡 Partial	Conceptual in docs, not implemented
Prediction drift	🔴 Missing	Not implemented
Performance monitoring	🟡 Partial	<code>evaluate_predictions.py</code> has metrics
Alerts/Visualizations	🔴 Missing	No Prometheus/Grafana

Deliverables:

- Evaluation metrics (accuracy, F1, confusion matrix)
- Confidence analysis in `evaluate_predictions.py`
- Drift detection implementation
- Prognosis model integration
- Monitoring dashboards

🔴 Month 5: Refinement & Retraining (10% Complete)

Task	Status	Evidence
Pipeline robustness	🟡 Partial	Error handling exists
Logging improvements	<input checked="" type="checkbox"/> Done	Structured logging in all scripts
Retraining trigger design	🔴 Missing	Not implemented
Retraining prototype	🔴 Missing	Not implemented

🔴 Month 6: Thesis Writing (5% Complete)

Task	Status	Evidence
Thesis document	🔴 Not Started	Only <code>Thesis_Plan.md</code> exists
Final presentation	🔴 Not Started	-
Code documentation	🟡 Partial	READMEs exist, docstrings in code
GitHub organization	<input checked="" type="checkbox"/> Done	Well-organized folder structure

📊 Component-Level Completion

pie title Pipeline Component Completion

"Data Ingestion" : 100
 "Preprocessing" : 100
 "Model Training" : 90
 "MLflow Tracking" : 95
 "DVC Versioning" : 100

"Docker Containers" : 85
"FastAPI Serving" : 80
"CI/CD Pipeline" : 10
"Drift Detection" : 15
"Monitoring" : 20
"Retraining" : 5
"Thesis Writing" : 5



Improvement Recommendations

Note: All recommendations below are backed by research papers from the literature review. Paper references are provided in [brackets] for verification.



Critical (Must Do)

1. Implement CI/CD Pipeline with GitHub Actions

Current State: No `.github/workflows/` directory

Recommendation: Create automated testing and deployment pipeline

Research Support:

- **[“Enabling End-To-End Machine Learning Replicability”]** - Emphasizes Docker and reproducibility for ML pipelines
- **[“Reproducible workflow for online AI in digital health”]** - CI/CD and versioning for healthcare applications
- **[“MLOps: A Survey”]** - Comprehensive MLOps lifecycle best practices
- **[“MLDev: Data Science Experiment Automation”]** - Experiment tracking automation with DVC, MLflow

```
# .github/workflows/ci.yml (recommended)
name: CI Pipeline
on: [push, pull_request]
jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - uses: actions/setup-python@v4
        with:
          python-version: '3.11'
      - run: pip install -r config/requirements.txt
      - run: pytest tests/
      - run: pylint tests/ -v
      - lint:
          runs-on: ubuntu-latest
          steps:
            - uses: actions/checkout@v3
            - run: pip install pylint
            - run: pylint src/ --rcfile=config/.pylintrc
      docker:
        runs-on: ubuntu-latest
        steps:
          - uses: actions/checkout@v3
          - run: docker build -f docker/Dockerfile.inference -t har-inference .
```

2. Implement Drift Detection

Current State: Only conceptual in documentation

Recommendation: Add statistical drift detection to `data_validator.py`

Research Support:

- **["Domain Adaptation for IMU-based HAR: A Survey"]** - Documents 40%+ accuracy drop due to domain shift
- **["Recognition of Anxiety-Related Activities using 1DCNN-BiLSTM" (ICTH_16)]** - Core paper showing lab-to-life gap (49% → 87% with adaptation)
- **["Are Anxiety Detection Models Generalizable"]** - Cross-population and cross-device studies on model degradation
- **["Resilience of ML Models in Anxiety Detection"]** - Impact of sensor noise on model performance

```
# Add to src/data_validator.py
from scipy.stats import ks_2samp

def detect_data_drift(reference_data, production_data, threshold=0.05):
    """ Detect data drift using Kolmogorov-Smirnov test. Args: reference_data: Training data statistics production_data: New production data threshold: P-value threshold for drift detection Returns: dict: Drift detection results per sensor """
    sensors = ['Ax', 'Ay', 'Az', 'Gx', 'Gy', 'Gz']
    drift_results = {}
    for idx, sensor in enumerate(sensors):
        ref_flat = reference_data[:, :, idx].flatten()
        prod_flat = production_data[:, :, idx].flatten()
        stat, p_value = ks_2samp(ref_flat, prod_flat)
        drift_results[sensor] = {
            'ks_statistic': float(stat),
            'p_value': float(p_value),
            'drift_detected': p_value < threshold
        }
    return drift_results
```

3. Add Unit Tests

Current State: `tests/` folder is empty

Recommendation: Create test files for critical components

Research Support:

- **["Toward Reusable Science with Readable Code"]** - Best practices for reproducible scientific code
- **["MLOps: A Survey"]** - Testing as critical component of ML lifecycle
- **["DevOps-Driven Real-Time Health Analytics"]** - DevOps practices including testing for health pipelines

```
# tests/test_preprocessing.py (recommended)
import pytest
import numpy as np
from src.sensor_data_pipeline import SensorDataPipeline
```

```

def test_windowing():
    """Test sliding window creation."""
    data = np.random.randn(1000, 6)
    windows = create_windows(data, window_size=200, overlap=0.5)
    assert windows.shape[1] == 200  assert windows.shape[2] == 6
def test_unit_conversion():
    """Test milliG to m/s2 conversion."""
    milli_g = 1000 # 1G in milliG  m_s2 = milli_g * 0.00981  assert abs(m_s2 - 9.81) < 0.01
def test_domain_calibration():
    """Test domain calibration offset."""
    offset = -6.295  raw_az = -9.81 # Raw gravity  calibrated = raw_az + offset
    assert abs(calibrated - (-3.515)) < 0.1

```

🟡 Important (Should Do)

4. Add Self-Attention Layer to Model

Current State: 1D-CNN + BiLSTM only

Recommendation: Add attention for improved temporal modeling

Research Support:

- **["Deep CNN-LSTM With Self-Attention Model for Human Activity Recognition"]** - Hybrid CNN-LSTM with self-attention achieves +2-3% accuracy improvement
- **["CNNs, RNNs and Transformers in HAR: A Survey"]** - Comparison showing attention benefits for long-range temporal dependencies
- **["Human Activity Recognition using Multi-Head CNN followed by LSTM"]** - Multi-head architecture for enhanced feature extraction

```

# Enhancement for model architecture
from tensorflow.keras.layers import MultiHeadAttention, Add, LayerNormalization
class AttentionBlock(tf.keras.layers.Layer):
    def __init__(self, num_heads=4, key_dim=64):
        super().__init__()
        self.attention = MultiHeadAttention(num_heads=num_heads, key_dim=key_dim)
        self.norm = LayerNormalization()
    def call(self, x):
        attention_output = self.attention(x, x)
        return self.norm(x + attention_output)

```

Expected Improvement: +1-3% accuracy based on research papers

5. Implement Data Augmentation

Current State: No augmentation applied

Recommendation: Add jittering, scaling, time-warping

Research Support:

- **["Deep learning for sensor-based activity recognition: A survey"]** - Documents standard augmentation techniques for HAR
- **["Self-supervised learning for HAR"]** - Self-supervised pre-training with augmentation strategies
- **["Masked Video and Body-worn IMU Autoencoder"]** - Masked autoencoder approach with data augmentation
- **["Comparative Study on the Effects of Noise in HAR"]** - Understanding noise tolerance for robust training

```
# src/data_augmentation.py (new file)import numpy as np
def augment_sensor_data(X, augmentation_config):
    """Apply data augmentation to sensor windows."""
    augmented = X.copy()
    # Jittering: Add Gaussian noise if augmentation_config.get('jitter', False):
        noise = np.random.normal(0, 0.05, X.shape)
        augmented += noise
    # Scaling: Random magnitude scaling if augmentation_config.get('scale', False):
        scale = np.random.uniform(0.9, 1.1)
        augmented *= scale
    return augmented
```

6. Add 5-Fold Cross-Validation

Current State: Single train/test split

Recommendation: Implement k-fold CV for robust evaluation

Research Support:

- **["Recognition of Anxiety-Related Activities using 1DCNN-BiLSTM" (ICTH_16)]** - Uses 5-fold stratified CV as standard methodology
- **["A Close Look into Human Activity Recognition Models using Deep Learning"]** - Comprehensive comparison using k-fold validation
- **["Transfer Learning in HAR: A Survey"]** - Cross-validation for domain adaptation evaluation

```
# Add to evaluate_predictions.pyfrom sklearn.model_selection import StratifiedKFold
def cross_validate_model(model, X, y, n_splits=5):
    """Perform stratified k-fold cross-validation."""
    skf = StratifiedKFold(n_splits=n_splits, shuffle=True, random_state=42)
    fold_metrics = []
    for fold, (train_idx, val_idx) in enumerate(skf.split(X, y)):
        X_train, X_val = X[train_idx], X[val_idx]
        y_train, y_val = y[train_idx], y[val_idx]
        # Train and evaluate
        model.fit(X_train, y_train, epochs=50, verbose=0)
```

```

y_pred = model.predict(X_val).argmax(axis=1)
f1 = f1_score(y_val, y_pred, average='macro')
fold_metrics.append({'fold': fold, 'f1_macro': f1})
return fold_metrics

```

Nice to Have (Optional)

7. RAG-Enhanced Clinical Reports

Recommendation: Add LLM-based report generation for clinical insights

Research Support:

- **["A Multi-Stage, RAG-Enhanced Pipeline for Generating Mental Health Reports" (EHB_2025_71)]** - Core architecture: HAR → bout analysis → LLM reports
- **["Retrieval-Augmented Generation (RAG) in Healthcare: A Comprehensive Review"]**
- RAG architecture patterns for healthcare
- **["Medical Graph RAG: Towards Safe Medical Large Language Models"]** - Knowledge graph-based RAG for medical safety
- **["Evaluating LLMs on medical evidence summarization"]** - Highlights hallucination risks mitigated by RAG
- **["Enhancing RAG with Knowledge Graph-Elicited Reasoning"]** - Knowledge graphs for healthcare copilots

```

# Future: src/report_generator.py class ClinicalReportGenerator:
    """Generate clinical reports from HAR predictions using RAG."""
    def __init__(self, knowledge_base, llm_model):
        self.kb = knowledge_base # Neo4j or vector store      self.llm = llm_model
    def generate_summary(self, predictions, bout_analysis):
        # Retrieve relevant clinical context      context = self.kb.retrieve(predictions.top
        _activities)
        # Generate report with LLM      prompt = f"Based on activity patterns: {bout_an
        alysis}..."      return self.llm.generate(prompt)

```

8. Temporal Bout Analysis

Recommendation: Aggregate consecutive predictions into "bouts"

Research Support:

- **["A Multi-Stage, RAG-Enhanced Pipeline for Generating Mental Health Reports" (EHB_2025_71)]** - Defines bout analysis with gap thresholds (120s HR, 300s behavior)
- **["A State-of-the-Art Review of Computational Models for Longitudinal Wearable Data"]** - Temporal aggregation methods
- **["Deep Learning Paired with Wearable Passive Sensing Data Predicts Anxiety Deterioration"]** - Longitudinal pattern analysis over 17-18 years

- [“Designing a Clinician-Centered Wearable Data Dashboard (CarePortal)"] - Clinical visualization of bout-level data

```
# src/bout_analysis.py (new file)
def detect_bouts(predictions, timestamps, min_duration_sec=5):
    """ Detect activity bouts (consecutive same-activity periods). Args:
        predictions: Array of predicted activity labels
        timestamps: Array of corresponding time stamps
        min_duration_sec: Minimum bout duration to consider
    Returns:
        DataFrame with bout statistics
    """
    bouts = []
    current_activity = predictions[0]
    bout_start = timestamps[0]
    for i, (pred, ts) in enumerate(zip(predictions[1:], timestamps[1:]), 1):
        if pred != current_activity:
            duration = (timestamps[i-1] - bout_start).total_seconds()
            if duration >= min_duration_sec:
                bouts.append({
                    'activity': current_activity,
                    'start': bout_start,
                    'end': timestamps[i-1],
                    'duration_sec': duration
                })
            current_activity = pred
            bout_start = ts
    return pd.DataFrame(bouts)
```

Priority Matrix

quadrantChart
 title Priority Matrix for Improvements
 x-axis Low Effort → High Effort
 y-axis Low Impact → High Impact
 quadrant-1 Do First
 quadrant-2 Plan Carefully
 quadrant-3 Quick Wins
 quadrant-4 Consider Later

CI/CD Pipeline: [0.7, 0.9]
 Drift Detection: [0.5, 0.85]
 Unit Tests: [0.3, 0.75]
 Self-Attention: [0.6, 0.5]
 Data Augmentation: [0.35, 0.55]
 Cross-Validation: [0.25, 0.6]

RAG Reports: [0.9, 0.4]
 Bout Analysis: [0.4, 0.35]

Recommended Action Plan

Week 1-2 (December 15-28, 2025)

Priority	Task	Effort	Impact
🔴	Create GitHub Actions CI/CD	Medium	High
🔴	Add unit tests for preprocessing	Low	High
🔴	Implement drift detection	Medium	High

Week 3-4 (December 29 - January 11, 2026)

Priority	Task	Effort	Impact
🟡	Add 5-fold cross-validation	Low	Medium
🟡	Implement data augmentation	Medium	Medium
🟡	Add F1-macro as primary metric	Low	Medium

Week 5-8 (January 12 - February 8, 2026)

Priority	Task	Effort	Impact
🟡	Self-attention layer	Medium	Medium
🟡	Retraining trigger prototype	High	High
🟡	Prometheus/Grafana monitoring	High	Medium

Week 9-12 (February 9 - March 8, 2026)

Priority	Task	Effort	Impact
🟢	Bout analysis	Medium	Low
🟢	RAG report generation	High	Low
📝	Thesis writing begins	High	Critical

Current Repository Strengths

Strength	Evidence
<input checked="" type="checkbox"/> Well-organized folder structure	Clear separation: <code>src/</code> , <code>data/</code> , <code>models/</code> , <code>config/</code> , <code>docs/</code>
<input checked="" type="checkbox"/> Comprehensive documentation	8+ markdown files in <code>docs/</code> , 77+ papers analyzed
<input checked="" type="checkbox"/> MLflow integration	<code>mlflow_tracking.py</code> (654 lines), experiments tracked

Strength	Evidence
✓ DVC data versioning	.dvc files for raw, processed, prepared, models
✓ Docker containerization	Training and inference Dockerfiles
✓ FastAPI serving	docker/api/main.py with endpoints
✓ Domain adaptation	Calibration offset applied (-6.295 m/s ²)
✓ Evaluation metrics	Accuracy, F1, confusion matrix, confidence analysis

Current Repository Gaps

Gap	Impact	Recommended Fix
✗ No CI/CD pipeline	High	Create .github/workflows/
✗ No unit tests	High	Add tests/ with pytest
✗ No drift detection	High	Add to data_validator.py
✗ No prognosis model	Medium	Design data flow
✗ No monitoring dashboard	Medium	Add Prometheus/Grafana
✗ No retraining trigger	Medium	Prototype trigger logic
✗ No data augmentation	Low	Add augmentation module
✗ No self-attention	Low	Modify model architecture

Success Metrics

Metric	Current	Target	Gap
Pipeline completion	58%	100%	42%
Model accuracy	87%	90%+	3%+
Test coverage	0%	80%	80%
CI/CD automation	0%	100%	100%
Documentation	70%	100%	30%
Drift detection	0%	100%	100%



Conclusion

The thesis is approximately **58% complete** with strong foundations in:

- Data ingestion and preprocessing
- Model training and MLflow tracking
- Docker containerization and API serving

Critical gaps requiring immediate attention:

1. CI/CD pipeline (GitHub Actions)

2. Unit tests
3. Drift detection implementation

Timeline assessment: On track if CI/CD and monitoring are completed by end of January 2026, leaving February-March for refinement and thesis writing.

Paper References (Full Citations)

All improvement recommendations are supported by the following research papers from your literature collection ([research_papers/76 papers/](#)):

Core Methodology Papers (★★★★★)

Citation Key	Full Title	Location
ICTH_16	"Recognition of Anxiety-Related Activities using 1DCNN-BiLSTM on Commercial Wearable"	research_papers/76 papers/ICTH_16/
EHB_2025_71	"A Multi-Stage, RAG-Enhanced Pipeline for Generating Mental Health Reports from Wearable Sensor Data"	research_papers/76 papers/EHB_2025_71/
ADAM-sense	Khan et al. (2021) - "ADAM-sense: Anxiety-displaying activities recognition by motion sensors"	research_papers/76 papers/

HAR & Deep Learning Papers

Paper Title	Key Contribution	Used For
"Deep CNN-LSTM With Self-Attention Model for Human Activity Recognition"	Self-attention for HAR	Self-attention recommendation
"A Close Look into Human Activity Recognition Models using Deep Learning"	BiLSTM, CNN comparison	Architecture choices
"Deep learning for sensor-based activity recognition: A survey"	Foundational DL for HAR	Data augmentation, windowing
"Human Activity Recognition using Multi-Head CNN followed by LSTM"	Multi-head architecture	Multi-scale feature extraction
"CNNs, RNNs and Transformers in HAR: A Survey and Hybrid Model"	Transformer comparison	Attention mechanisms
"Combining Accelerometer and Gyroscope Data in Smartphone-Based Activity Recognition"	Sensor fusion	Preprocessing validation

Mental Health & Anxiety Detection Papers

Paper Title	Key Contribution	Used For
"Wearable Artificial Intelligence for Detecting Anxiety: Systematic Review"	Meta-analysis of wearable AI	Overall approach validation

Paper Title	Key Contribution	Used For
"Are Anxiety Detection Models Generalizable: Cross-Activity and Cross-Population Study"	Generalization challenges	Drift detection rationale
"Resilience of ML Models in Anxiety Detection: Impact of Gaussian Noise"	Noise robustness	Data augmentation rationale
"Deep Learning Paired with Wearable Passive Sensing Data Predicts Anxiety Deterioration"	Longitudinal prediction	Bout analysis rationale
"Panic Attack Prediction Using Wearable Devices and Machine Learning"	Time-series prediction	Real-time inference

MLOps & Pipeline Papers

Paper Title	Key Contribution	Used For
"MACHINE LEARNING OPERATIONS: A SURVEY ON MLOPS"	MLOps lifecycle	CI/CD, testing recommendations
"MLHOps: Machine Learning for Healthcare Operations"	Healthcare MLOps	Compliance considerations
"Enabling End-To-End Machine Learning Replicability"	Reproducibility	Docker, CI/CD
"MLDEV: DATA SCIENCE EXPERIMENT AUTOMATION"	DVC, MLflow	Experiment tracking
"Reproducible workflow for online AI in digital health"	Digital health workflows	CI/CD pipeline
"DevOps-Driven Real-Time Health Analytics"	DevOps for health	Testing practices
"Toward Reusable Science with Readable Code"	Code quality	Unit testing rationale

Domain Adaptation Papers

Paper Title	Key Contribution	Used For
"Domain Adaptation for IMU-based HAR: A Survey"	Domain shift analysis	Drift detection
"Transfer Learning in HAR: A Survey"	Transfer learning	Cross-validation, fine-tuning

RAG & LLM Papers

Paper Title	Key Contribution	Used For
"Retrieval-Augmented Generation (RAG) in Healthcare: A Comprehensive Review"	RAG architecture	Clinical reports
"Medical Graph RAG: Towards Safe Medical Large Language Models"	Knowledge graph RAG	Safety considerations

Paper Title	Key Contribution	Used For
"Evaluating LLMs on medical evidence summarization"	Hallucination risks	RAG rationale
"Enhancing RAG with Knowledge Graph-Elicited Reasoning"	Knowledge graphs	Report generation

Foundation Models & Self-Supervised Learning Papers

Paper Title	Key Contribution	Used For
"Toward Foundation Model for Multivariate Wearable Sensing"	Foundation models	Future improvements
"Masked Video and Body-worn IMU Autoencoder"	MAE for IMU	Augmentation strategies
"Self-supervised learning for fast and scalable time series"	SSL for HPO	Training efficiency

Wearable Data Processing Papers

Paper Title	Key Contribution	Used For
"A State-of-the-Art Review of Computational Models for Longitudinal Wearable Data"	Temporal analysis	Bout analysis
"Designing a Clinician-Centered Wearable Data Dashboard (CarePortal)"	Clinical visualization	Dashboard recommendations
"Comparative Study on the Effects of Noise in HAR"	Noise analysis	Augmentation validation

Quick Reference Guide

When implementing improvements, refer to these papers:

Improvement	Primary Papers to Read
CI/CD Pipeline	"MLOps Survey", "End-to-End ML Replicability", "Reproducible Workflow"
Drift Detection	ICTH_16, "Domain Adaptation Survey", "Are Anxiety Models Generalizable"
Self-Attention	"Deep CNN-LSTM With Self-Attention", "CNNs/RNNs/Transformers Survey"
Data Augmentation	"Deep Learning for HAR Survey", "Masked IMU Autoencoder"
Cross-Validation	ICTH_16, "A Close Look into HAR Models"
Bout Analysis	EHB_2025_71, "Longitudinal Wearable Data Review"
RAG Reports	EHB_2025_71, "RAG in Healthcare Review", "Medical Graph RAG"

Generated: December 13, 2025

Based on: [Thesis_Plan.md](#), repository analysis, and 77+ research papers

Full paper summaries available in: [research_papers/COMPREHENSIVE_RESEARCH_PAPERS_SUMMARY.md](#)